

Project Report 3

The ms (“memory status”) Program

Team Members:
Howard Grimberg
Aditya Balasubramanian

Submitted: 5/9/2013

Purpose

The purpose of the ms program is to use the /proc virtual file system in Linux to read various statistics about the memory on the system. Cumulative totals and per process statistics were gathered and displayed in such a way to make comparisons easy.

Implementation

For the cumulative totals, the '/proc/meminfo' file was scanned and the choice bits extracted and formatted neatly. Not all statistics are available in that location, so other metrics had to be computed differently. For example, the amount of memory in use was calculated by subtracting the free memory from the total amount of memory available. The other cumulative metrics were tabulated by scanning the entire ProcessList object and summing the appropriate values.

For the individual totals, all PIDs (in use or not) were obtained by scanning the /proc directory and capturing everything that was an integer. Each PID was then given an instance of the Process class that was responsible for updating itself when asked (explained later). The Process class will scan the subdirectory of its PID in /proc and obtain the necessary info from the smaps and the cmdline file. The smaps file contains statistics for each chunk of memory allocated to that given PID. As such, a sum was taken across all of the memory blocks with some not-very-clever text processing kung-fu. Originally, I wanted the ms program to be more like top and be constantly updating rather than taking a snapshot. For that reason, every PID has its own class that updates itself. Sorting was done using the builtin in sorted() routine in python and careful massaging of the ProcessList class into tuples.

Testing

Testing was very simple. The project was tested on a personal linux computer (so permissions were not an issue). A random, large process was chosen (rhythmbox) and the output captured using ms. Top and gnome-system-monitor were used to confirm that the memory values were about the same. Memory maps can be obtained in gnome-system-monitor by right clicking on a process and selecting memory maps. The program was then run in the VM with the specified test subjects and documented. One terminal was used to run the test subject while the other terminal was used to run ms rapidly in succession (using fast keystrokes and bash history) after a 10 second warm-up time. The swap file was turned off and on using swapon -a and swapon -a between each test to ensure no residual data was left in swap space.

Process Description

Mystery Program Aether :

Total Usage

Total 297832 kB

In Use 66244 kB

Free 231588 kB

Shared 114748 kB

Nonshared 12044 kB

ID	USS	PSS	SWAP	RES	SHR	CMD
2509		24	10270	0	20636	20612 ./aether
2507		24	6856	0	20636	20612 ./aether
2506		24	6856	0	20636	20612 ./aether
2505		24	6856	0	20636	20612 ./aether
2508		24	10270	0	20636	20612 ./aether

The aether process follows the below memory characteristics .Processes 2509 and 2508 share the same page size of $(10270 * 2 = 20540 \text{ KB})$ and processes 2505,2506 and 2507 share the page size of $(6586 * 3) 19758 \text{ KB}$.

- (d) Two buffers of the same size are allocated. Two children share one of the buffers, and the other three children share the other. Each child requires only a small amount of unique memory.

Mystery Program hemera - A

ID	USS	PSS	SWAP	RES	SHR	CMD
2524		20504	20510	0	20644	140 ./hemera
2523		16408	16414	0	16548	140 ./hemera

2522	12312	12318	0	12452	140 ./hemera
2521	8216	8222	0	8356	140 ./hemera
2520	4120	4126	0	4260	140 ./hemera

The hemera process follows the below memory characteristic. The Unique memory is relatively large as compared to the Shared memory.

- (a) each child allocates their own (relatively large) buffer of unique memory. Each child requires a different amount of unique memory. Each child shares a relatively small portion of its address space with the other children.

Mystery Program ourea-C

Total Usage

Total 297832 kB

In Use 45844 kB

Free 251988 kB

Shared 114964 kB

Nonshared 11936 kB

PID	USS	PSS	SWAP	RES	SHR	CMD
2549	24	4124	0	20628	20604	./ourea
2551	24	4124	0	20628	20604	./ourea
2552	24	4124	0	20628	20604	./ourea
2550	24	4124	0	20628	20604	./ourea
2553	24	4124	0	20628	20604	./ourea

The Mystery program ourea follows the statement C characteristic as each process has a small unique memory (USS 24) and the Shared memory buffer is large and same for all the processes which is SHR 20604

- (c) The children all share a single, large common buffer. Each child requires only a small amount of unique memory.

Mystery Program pontus-E

Total Usage

Total 297832 kB

In Use 86884 kB

Free 210948 kB

Shared 73996 kB

Nonshared 52892 kB

PID	USS	PSS	SWAP	RES	SHR	CMD
2564	4120	9586	0	20628	16508	./pontus
2561	16408	17778	0	20628	4220	./pontus
2563	8216	12316	0	20628	12412	./pontus
2562	12312	15047	0	20628	8316	./pontus
2565	24	6855	0	20620	20596	./pontus

The Mystery Program Pontus follows the below characteristic.

As can be seen each process shares the same amount of Resident memory (RES) but has different values for Shared memory which it has shared with other children (SHR).

- (e) Each child requires about the same amount of resident memory. However, each child shares a different amount of its required memory with the other children.

Mystery Program Uranus-F

Total Usage

Total 297832 kB

In Use 128404 kB

Free 169428 kB

Shared 19056 kB

Nonshared 114740 kB

PID	USS	PSS	SWAP	RES	SHR	CMD
2570	20600	20775	0	22008	1408	./uranus
2569	20600	20775	0	22008	1408	./uranus
2571	20600	20775	0	22008	1408	./uranus
2572	20600	20775	0	22008	1408	./uranus
2573	20600	20775	0	22008	1408	./uranus

Mystery Program Uranus follows the F description as the size of unique buffer is same for all the processes and it the Shared memory is small relatively to the amount of unique memory. The difference between this process and the other processes is that it doesn't swap out its memory when it is out of physical memory.

We ran the process with limited memory of 300MB and ran the same process in other terminals (before all the processes completed) such that the entire 300 MB RAM is exhausted. We continuously monitored the memory status by running `ms` each time. After a while when we tried to run the process again, we couldn't because it had run out of memory and the memory wasn't swapped out to the disk.

- (f) Each child allocates their own (relatively large) buffer of unique memory. Each child's unique buffer is about the same size as each other child. Each child shares a relatively small portion of its address space with the other children. In contrast to all of the other programs, this program uses a system call to lock all of its allocated memory in RAM. None of its allocated memory will be swapped out to disk, even when the system runs out of memory

Mystery Program typhoon-B

Total Usage

Total 297832 kB

In Use 127092 kB

Free 170740 kB

Shared 12328 kB

Nonshared 114304 kB

PID	USS	PSS	SWAP	RES	SHR	CMD
2483	20504	20509	0	20632	128	./typhon
2486	20504	20509	0	20632	128	./typhon
2484	20504	20509	0	20632	128	./typhon
2485	20504	20509	0	20632	128	./typhon
2482	20504	20509	0	20632	128	./typhon

The Mystery Program typhon follows the below statement characteristic as the unique memory for each memory is same and large in this case (USS is 20504) and Shared memory is small (SHR is 128).

- (b) Each child allocates their own (relatively large) buffer of unique memory. Each child's unique buffer is about the same size as each other child. Each child shares a relatively small portion of its address space with the other children.